




Docker compose



Why we need orchestration

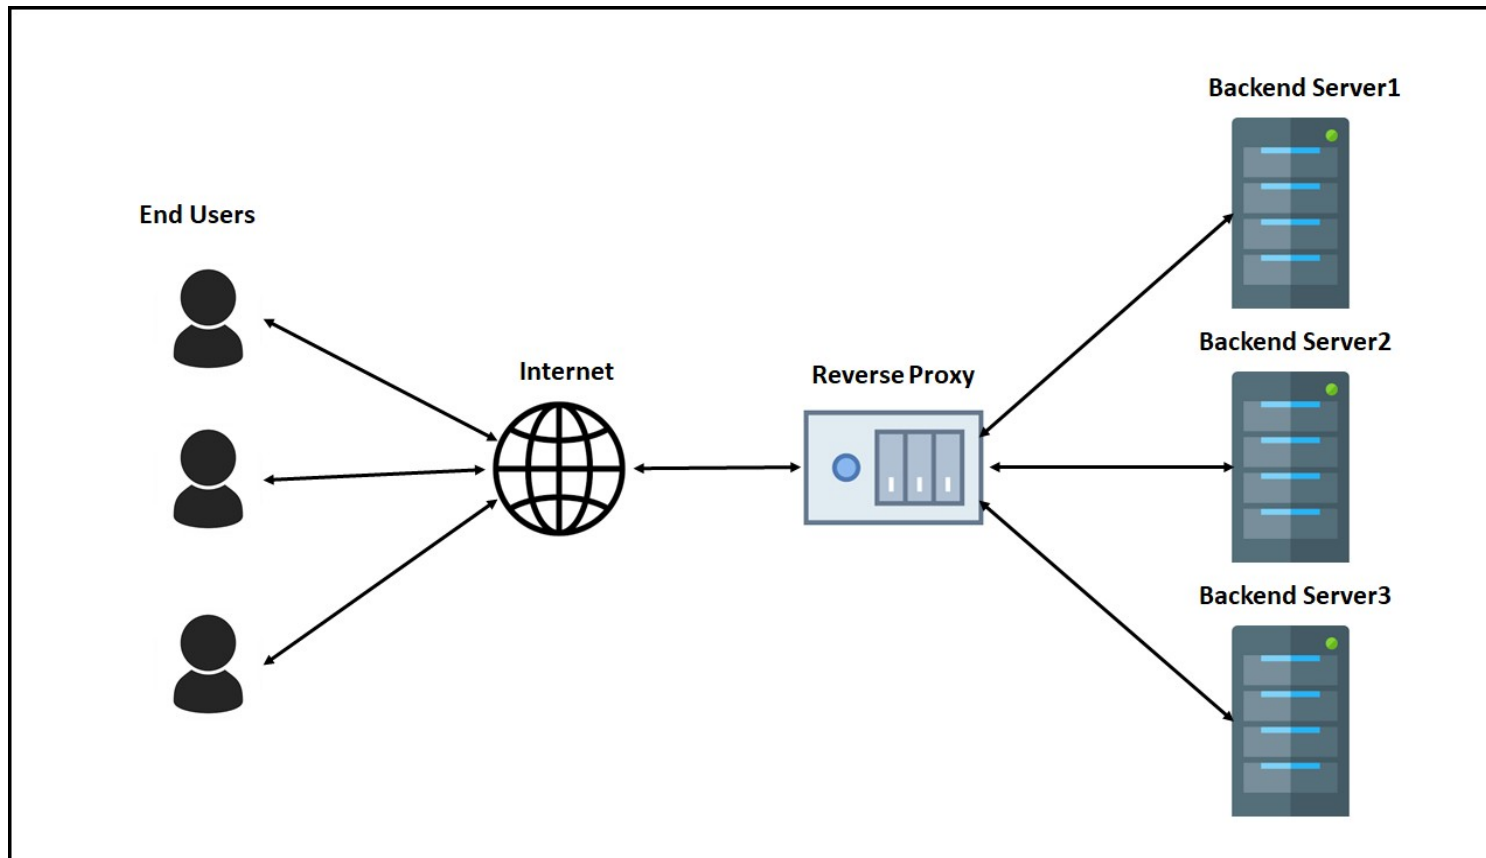
- Run an ecosystem of applications
- Prevent downtime during deployments
- Networking and routing capabilities
- Scalability capabilities

Deployments downtime



Networking / routing

- K8s offer a wide range of routing possibilities
es. Nginx, traefik, calico, istio

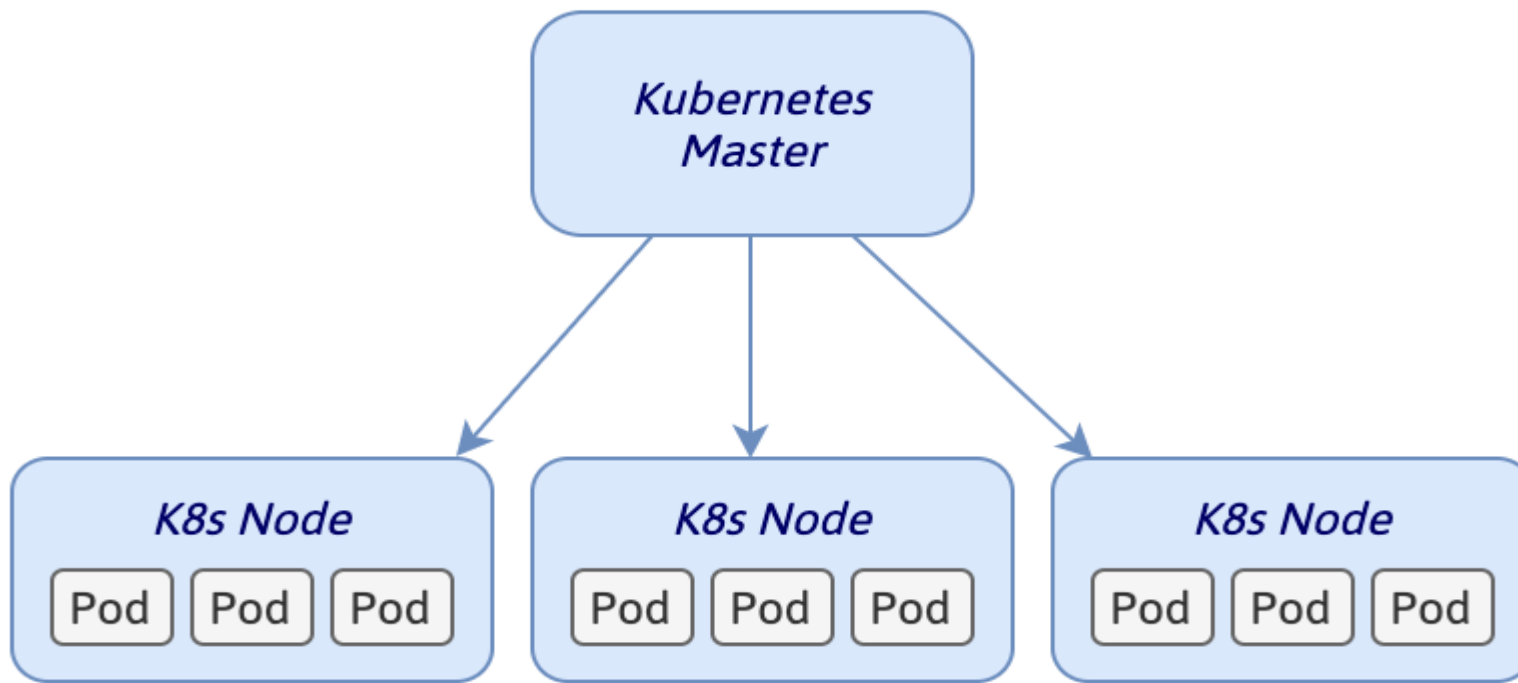


Scalability

- K8s clusters are composed by nodes. A pool of resources (CPU / RAM) is then shared across pods. A Pod is a set of containers



Kubernetes Cluster





Orchestration basics

- **Docker-compose**, docker swarm
- Kubernetes, minikube, microk8s
- Openshift (Paas)
- Rancher (hybrid cloud)
- Others: Mesos, Nomad

Orchestration basics

Docker compose has limited orchestration capabilities:


- use **depends_on** to define start priorities
- **Restart policy**
- **volumes** and **network**
- Basically is a transposition of the docker syntax to a yaml file

Volumes declaration syntax

- **host-mounted** volumes
/host/path:/container/path
- **named** volumes
volume_name:/container/path
- **shared** volumes
--volumes-from container_name

Volumes mount types

- Volumes are stored in a part of the host filesystem which is managed by Docker (/var/lib/docker/volumes/ on Linux). Non-Docker processes should not modify this part of the filesystem. Volumes are the best way to persist data in Docker.
- Bind mounts may be stored anywhere on the host system. They may even be important system files or directories. Non-Docker processes on the Docker host or a Docker container can modify them at any time.
- tmpfs mounts are stored in the host system's memory only, and are never written to the host system's filesystem.



Networking

```
netstat -rn      aureliano@aureliano-N141CU ~/git/devops-B3-2023 (main)$ netstat -rn
ifconfig        Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	MSS	Window	irrt	Iface
0.0.0.0	192.168.1.254	0.0.0.0	UG	0 0	0	wlp0s20f3	
10.0.3.0	0.0.0.0	255.255.255.0	U	0 0	0	lxcbr0	
169.254.0.0	0.0.0.0	255.255.0.0	U	0 0	0	wlp0s20f3	
172.17.0.0	0.0.0.0	255.255.0.0	U	0 0	0	docker0	

Can customize the ip range of docker.

```
Docker network ls
Docker network inspect my_network
Cat /etc/docker/daemon.json
```

Frequent use case: I want to connect from a container to a service on the host

The host has a changing IP address, or none if you have no network access. We recommend that you connect to the special DNS name `host.docker.internal`, which resolves to the internal IP address used by the host.

You can also reach the gateway using `gateway.docker.internal` (**By default, Docker uses 172.17.0.0/16.**)

Network declaration syntax

Can create subnets. This enforce isolation between containers

```
version: '3.6'
services:
  api-gateway:
    container_name: api-gateway
    image: api-gateway
    networks:
      - gateway
    ports:
      - 9090:8080
    restart:
      on-failure
  api-gateway-replica:
    container_name: api-gateway-
    replica
    image: api-gateway
    networks:
      - gateway-replica
    ports:
      - 9092:8080
    restart:
      on-failure
networks:
  gateway: {}
  gateway-replica: {}
```

Can customize the ip range of docker.

```
Docker network ls
Docker network inspect my_network
Cat /etc/docker/daemon.json
```

Orchestration - docker compose

Compose is a tool for defining and running multi-container Docker applications. With Compose, you use a YAML file to configure your application's services. Then, with a single command, you create and start all the services from your configuration.

It also has commands for managing the whole lifecycle of your application:

- Start, stop, and rebuild services
- View the status of running services
- Stream the log output of running services
- Run a one-off command on a service

The key features of Compose that make it effective are:

- Have multiple isolated environments on a single host
- Preserve volume data when containers are created
- Only recreate containers that have changed
- Support variables and moving a composition between environments

Orchestration - docker compose

version: "3.8"

services:

redis:

image: redis:7-bullseye
container_name: redis-airflow
restart: always
command: ["--databases", "1"]
healthcheck:
test: ["CMD", "redis-cli", "ping"]
interval: 10s
timeout: 5s
retries: 5

ports:

- 6379:6379

volumes:

- redisdata:/data

*Start postgres and redis: **docker-compose up***

Expose container port to host port

postgres:

image: postgres:14.7-bullseye
container_name: postgres
volumes:
- pgdata_mlflow:/var/lib/postgresql/data/pg_data
- ./docker/postgresql/scripts:/var/lib/postgresql/scripts

restart: always

environment:

- POSTGRES_PORT=5432
- POSTGRES_PASSWORD=ynov
- POSTGRES_USER=ynovcampus
- POSTGRES_DB=ynov
- PGDATA=/var/lib/postgresql/data/pg_data

Host mounted volume declaration

volumes:

Env var declaration

pgvol:

backup_pg_airflow:

2023/24

Conteneurisation et Orchestration
de conteneurs - Aureliano Sinatra
Named volumes declaration

13 / 16

Orchestration - docker compose

version: "3.8"

services:

redis:

build:

Define the root (context) from which files are resolved by dockerfile

context: ./my/context

dockerfile: ./my/Dockerfile *Define the dockerfile if not at same level of docker-compose.yaml*

container_name: redis-airflow

restart: always

command: ["--databases", "1"]

healthcheck:

test: ["CMD", "redis-cli", "ping"]

interval: 10s

timeout: 5s

retries: 5

...

build redis: **docker-compose build redis**



ENVIRONMENT VARIABLES

- DOTENV: env variables can be moved to .env files
- Another way is to set env var as SYSTEM or USER variables
- Set default ENV in the Dockerfile
- Override ENV at runtime via e flag or in the dockerfile



DOCKER MAINTENANCE

- Inspect: `ls /var/lib/docker/containers`
- `docker image prune`
- `docker container prune`
- `docker network prune`
- `docker volume prune`
- `docker system prune`